Model accuracy analysis

Portable metric summary. Open model_analysis.html for all SVG charts and full details.

model_name: accuracy_log_structured
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 255.114
regression.rmse_days: 327.303
regression.median_absolute_error_days: 187.372
regression.r2: -3.3237
regression.mean_error_days: 210.737
classification.accuracy: 0.972897
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.486448
classification.macro_f1: 0.493131
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0603861
classification.log_loss: 0.252657

model_name: accuracy_log_structured
split: test
split_role: test
n: 8299
regression.mae_days: 244.26
regression.rmse_days: 313.711
regression.median_absolute_error_days: 187.862
regression.r2: 0.00980404
regression.mean_error_days: 59.1994
classification.accuracy: 0.94421
classification.balanced_accuracy: 0.511365

classification.macro_precision: 0.513484
classification.macro_recall: 0.511365
classification.macro_f1: 0.512214
classification.precision: 0.058296
classification.recall: 0.0488722
classification.f1: 0.0531697
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.703692
classification.average_precision: 0.0673761
classification.brier: 0.0762873
classification.log_loss: 0.293566

model_name: accuracy_log_structured
split: validation
split_role: tuning
n: 4624
regression.mae_days: 251.851
regression.rmse_days: 322.667
regression.median_absolute_error_days: 195.593
regression.r2: 0.02169
regression.mean_error_days: 11.9989
classification.accuracy: 0.938365
classification.balanced_accuracy: 0.529086
classification.macro_precision: 0.535191
classification.macro_recall: 0.529086
classification.macro_f1: 0.531667
classification.precision: 0.105634
classification.recall: 0.0867052
classification.f1: 0.0952381
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.687022
classification.average_precision: 0.0730306
classification.brier: 0.0822933
classification.log_loss: 0.308743

model_name: accuracy_mean
split: calibration

split_role: calibration
n: 2841
regression.mae_days: 587.047
regression.rmse_days: 607.784
regression.median_absolute_error_days: 618.706
regression.r2: -13.9092
regression.mean_error_days: 587.047
classification.accuracy: 1
classification.balanced_accuracy: unavailable
classification.macro_precision: 1
classification.macro_recall: 1
classification.macro_f1: 1
classification.precision: unavailable
classification.recall: unavailable
classification.f1: unavailable
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0624149
classification.log_loss: 0.287455

model_name: accuracy_mean
split: test
split_role: test
n: 8299
regression.mae_days: 464.136
regression.rmse_days: 519.639
regression.median_absolute_error_days: 497.706
regression.r2: -1.71685
regression.mean_error_days: 413.081
classification.accuracy: 0.967948
classification.balanced_accuracy: 0.5
classification.macro_precision: 0.483974
classification.macro_recall: 0.5
classification.macro_f1: 0.491856
classification.precision: unavailable
classification.recall: 0

classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.5
classification.average_precision: 0.0320521
classification.brier: 0.0784518
classification.log_loss: 0.322697

model_name: accuracy_mean
split: validation
split_role: tuning
n: 4624
regression.mae_days: 412.349
regression.rmse_days: 478.323
regression.median_absolute_error_days: 417.706
regression.r2: -1.14986
regression.mean_error_days: 349.815
classification.accuracy: 0.962587
classification.balanced_accuracy: 0.5
classification.macro_precision: 0.481293
classification.macro_recall: 0.5
classification.macro_f1: 0.490468
classification.precision: unavailable
classification.recall: 0
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.5
classification.average_precision: 0.0374135
classification.brier: 0.0811344
classification.log_loss: 0.328592

model_name: accuracy_phase_shrinkage
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 508.276
regression.rmse_days: 572.911
regression.median_absolute_error_days: 436.894

regression.r2: -12.2474
regression.mean_error_days: 504.884
classification.accuracy: 0.965153
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.482577
classification.macro_f1: 0.491134
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0604679
classification.log_loss: 0.261094

model_name: accuracy_phase_shrinkage
split: test
split_role: test
n: 8299
regression.mae_days: 413.891
regression.rmse_days: 483.553
regression.median_absolute_error_days: 388.894
regression.r2: -1.35262
regression.mean_error_days: 357.139
classification.accuracy: 0.93698
classification.balanced_accuracy: 0.531257
classification.macro_precision: 0.527052
classification.macro_recall: 0.531257
classification.macro_f1: 0.528897
classification.precision: 0.0841424
classification.recall: 0.0977444
classification.f1: 0.0904348
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.700947
classification.average_precision: 0.0590582

classification.brier: 0.0769042
classification.log_loss: 0.300697

model_name: accuracy_phase_shrinkage
split: validation
split_role: tuning
n: 4624
regression.mae_days: 394.48
regression.rmse_days: 468.366
regression.median_absolute_error_days: 374.944
regression.r2: -1.06128
regression.mean_error_days: 322.406
classification.accuracy: 0.924092
classification.balanced_accuracy: 0.52445
classification.macro_precision: 0.520311
classification.macro_recall: 0.52445
classification.macro_f1: 0.521979
classification.precision: 0.0761905
classification.recall: 0.0924855
classification.f1: 0.0835509
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.634872
classification.average_precision: 0.0550867
classification.brier: 0.0852432
classification.log_loss: 0.322529

model_name: accuracy_task_selected
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 255.114
regression.rmse_days: 327.303
regression.median_absolute_error_days: 187.372
regression.r2: -3.3237
regression.mean_error_days: 210.737
classification.accuracy: 0.986624
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5

classification.macro_recall: 0.493312
classification.macro_f1: 0.496634
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0292237
classification.log_loss: 0.136724

model_name: accuracy_task_selected
split: test
split_role: test
n: 8299
regression.mae_days: 244.26
regression.rmse_days: 313.711
regression.median_absolute_error_days: 187.862
regression.r2: 0.00980404
regression.mean_error_days: 59.1994
classification.accuracy: 0.943728
classification.balanced_accuracy: 0.587449
classification.macro_precision: 0.575217
classification.macro_recall: 0.587449
classification.macro_f1: 0.580746
classification.precision: 0.176849
classification.recall: 0.206767
classification.f1: 0.190641
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.815134
classification.average_precision: 0.115926
classification.brier: 0.0525668
classification.log_loss: 0.2051

model_name: accuracy_task_selected
split: validation
split_role: tuning

n: 4624
regression.mae_days: 251.851
regression.rmse_days: 322.667
regression.median_absolute_error_days: 195.593
regression.r2: 0.02169
regression.mean_error_days: 11.9989
classification.accuracy: 0.928849
classification.balanced_accuracy: 0.613034
classification.macro_precision: 0.579597
classification.macro_recall: 0.613034
classification.macro_f1: 0.592471
classification.precision: 0.188
classification.recall: 0.271676
classification.f1: 0.222222
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.802266
classification.average_precision: 0.15468
classification.brier: 0.0618927
classification.log_loss: 0.227649

model_name: design_catboost
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 271.674
regression.rmse_days: 397.612
regression.median_absolute_error_days: 169.094
regression.r2: -5.38079
regression.mean_error_days: 214.041
classification.accuracy: 0.938402
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.469201
classification.macro_f1: 0.484111
classification.precision: 0
classification.recall: unavailable
classification.f1: 0

classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0527454
classification.log_loss: 0.196707

model_name: design_catboost
split: test
split_role: test
n: 8299
regression.mae_days: 242.41
regression.rmse_days: 336.062
regression.median_absolute_error_days: 174.536
regression.r2: -0.136324
regression.mean_error_days: 98.1035
classification.accuracy: 0.916014
classification.balanced_accuracy: 0.673093
classification.macro_precision: 0.574287
classification.macro_recall: 0.673093
classification.macro_f1: 0.597732
classification.precision: 0.168971
classification.recall: 0.413534
classification.f1: 0.239913
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.82044
classification.average_precision: 0.134345
classification.brier: 0.0719318
classification.log_loss: 0.251953

model_name: design_catboost
split: validation
split_role: tuning
n: 4624
regression.mae_days: 254.442
regression.rmse_days: 344.019
regression.median_absolute_error_days: 192.95
regression.r2: -0.112071

regression.mean_error_days: 65.9107
classification.accuracy: 0.908088
classification.balanced_accuracy: 0.621695
classification.macro_precision: 0.561046
classification.macro_recall: 0.621695
classification.macro_f1: 0.57693
classification.precision: 0.15
classification.recall: 0.312139
classification.f1: 0.202627
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.773392
classification.average_precision: 0.137818
classification.brier: 0.0791573
classification.log_loss: 0.274671

model_name: design_linear
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 457.574
regression.rmse_days: 567.068
regression.median_absolute_error_days: 393.123
regression.r2: -11.9786
regression.mean_error_days: 416.938
classification.accuracy: 0.936994
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.468497
classification.macro_f1: 0.483736
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0600687

classification.log_loss: 0.229261

model_name: design_linear
split: test
split_role: test
n: 8299
regression.mae_days: 391.998
regression.rmse_days: 493.773
regression.median_absolute_error_days: 336.779
regression.r2: -1.45311
regression.mean_error_days: 305.591
classification.accuracy: 0.913484
classification.balanced_accuracy: 0.63362
classification.macro_precision: 0.559095
classification.macro_recall: 0.63362
classification.macro_f1: 0.576467
classification.precision: 0.14127
classification.recall: 0.334586
classification.f1: 0.198661
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.800465
classification.average_precision: 0.120074
classification.brier: 0.078014
classification.log_loss: 0.279085

model_name: design_linear
split: validation
split_role: tuning
n: 4624
regression.mae_days: 373.987
regression.rmse_days: 475.065
regression.median_absolute_error_days: 313.828
regression.r2: -1.12067
regression.mean_error_days: 267.522
classification.accuracy: 0.901168
classification.balanced_accuracy: 0.609767
classification.macro_precision: 0.551669
classification.macro_recall: 0.609767

classification.macro_f1: 0.564937
classification.precision: 0.132124
classification.recall: 0.294798
classification.f1: 0.182469
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.770598
classification.average_precision: 0.126288
classification.brier: 0.0857907
classification.log_loss: 0.298394

model_name: design_log_linear
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 257.639
regression.rmse_days: 362.593
regression.median_absolute_error_days: 177.608
regression.r2: -4.30634
regression.mean_error_days: 204.137
classification.accuracy: 0.936994
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.468497
classification.macro_f1: 0.483736
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0600687
classification.log_loss: 0.229261

model_name: design_log_linear
split: test
split_role: test
n: 8299

regression.mae_days: 248.863
regression.rmse_days: 344.372
regression.median_absolute_error_days: 183.238
regression.r2: -0.193216
regression.mean_error_days: 79.7079
classification.accuracy: 0.913484
classification.balanced_accuracy: 0.63362
classification.macro_precision: 0.559095
classification.macro_recall: 0.63362
classification.macro_f1: 0.576467
classification.precision: 0.14127
classification.recall: 0.334586
classification.f1: 0.198661
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.800465
classification.average_precision: 0.120074
classification.brier: 0.078014
classification.log_loss: 0.279085

model_name: design_log_linear
split: validation
split_role: tuning
n: 4624
regression.mae_days: 255.463
regression.rmse_days: 345.97
regression.median_absolute_error_days: 191.578
regression.r2: -0.124718
regression.mean_error_days: 39.6958
classification.accuracy: 0.901168
classification.balanced_accuracy: 0.609767
classification.macro_precision: 0.551669
classification.macro_recall: 0.609767
classification.macro_f1: 0.564937
classification.precision: 0.132124
classification.recall: 0.294798
classification.f1: 0.182469
classification.threshold: 0.5

classification.positive_count: 173
classification.auroc: 0.770598
classification.average_precision: 0.126288
classification.brier: 0.0857907
classification.log_loss: 0.298394

model_name: design_task_selected
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 255.114
regression.rmse_days: 327.303
regression.median_absolute_error_days: 187.372
regression.r2: -3.3237
regression.mean_error_days: 210.737
classification.accuracy: 0.986624
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.493312
classification.macro_f1: 0.496634
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0292237
classification.log_loss: 0.136724

model_name: design_task_selected
split: test
split_role: test
n: 8299
regression.mae_days: 244.26
regression.rmse_days: 313.711
regression.median_absolute_error_days: 187.862
regression.r2: 0.00980404
regression.mean_error_days: 59.1994

classification.accuracy: 0.943728
classification.balanced_accuracy: 0.587449
classification.macro_precision: 0.575217
classification.macro_recall: 0.587449
classification.macro_f1: 0.580746
classification.precision: 0.176849
classification.recall: 0.206767
classification.f1: 0.190641
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.815134
classification.average_precision: 0.115926
classification.brier: 0.0525668
classification.log_loss: 0.2051

model_name: design_task_selected
split: validation
split_role: tuning
n: 4624
regression.mae_days: 251.851
regression.rmse_days: 322.667
regression.median_absolute_error_days: 195.593
regression.r2: 0.02169
regression.mean_error_days: 11.9989
classification.accuracy: 0.928849
classification.balanced_accuracy: 0.613034
classification.macro_precision: 0.579597
classification.macro_recall: 0.613034
classification.macro_f1: 0.592471
classification.precision: 0.188
classification.recall: 0.271676
classification.f1: 0.222222
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.802266
classification.average_precision: 0.15468
classification.brier: 0.0618927
classification.log_loss: 0.227649

model_name: median_baseline
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 286.214
regression.rmse_days: 318.896
regression.median_absolute_error_days: 309
regression.r2: -3.10444
regression.mean_error_days: 277.341
classification.accuracy: 1
classification.balanced_accuracy: unavailable
classification.macro_precision: 1
classification.macro_recall: 1
classification.macro_f1: 1
classification.precision: unavailable
classification.recall: unavailable
classification.f1: unavailable
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0624149
classification.log_loss: 0.287455

model_name: median_baseline
split: test
split_role: test
n: 8299
regression.mae_days: 285.51
regression.rmse_days: 331.776
regression.median_absolute_error_days: 288
regression.r2: -0.107521
regression.mean_error_days: 103.375
classification.accuracy: 0.967948
classification.balanced_accuracy: 0.5
classification.macro_precision: 0.483974
classification.macro_recall: 0.5
classification.macro_f1: 0.491856

classification.precision: unavailable
classification.recall: 0
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.5
classification.average_precision: 0.0320521
classification.brier: 0.0784518
classification.log_loss: 0.322697

model_name: median_baseline
split: validation
split_role: tuning
n: 4624
regression.mae_days: 281.284
regression.rmse_days: 328.681
regression.median_absolute_error_days: 271
regression.r2: -0.0151163
regression.mean_error_days: 40.1088
classification.accuracy: 0.962587
classification.balanced_accuracy: 0.5
classification.macro_precision: 0.481293
classification.macro_recall: 0.5
classification.macro_f1: 0.490468
classification.precision: unavailable
classification.recall: 0
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.5
classification.average_precision: 0.0374135
classification.brier: 0.0811344
classification.log_loss: 0.328592

model_name: multimodal
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 319.077

regression.rmse_days: 417.17
regression.median_absolute_error_days: 239.156
regression.r2: -6.02397
regression.mean_error_days: 276.354
classification.accuracy: 0.986624
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.493312
classification.macro_f1: 0.496634
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0292237
classification.log_loss: 0.136724

model_name: multimodal
split: test
split_role: test
n: 8299
regression.mae_days: 275.285
regression.rmse_days: 367.763
regression.median_absolute_error_days: 215.94
regression.r2: -0.360811
regression.mean_error_days: 172.43
classification.accuracy: 0.943728
classification.balanced_accuracy: 0.587449
classification.macro_precision: 0.575217
classification.macro_recall: 0.587449
classification.macro_f1: 0.580746
classification.precision: 0.176849
classification.recall: 0.206767
classification.f1: 0.190641
classification.threshold: 0.5
classification.positive_count: 266

classification.auroc: 0.815134
classification.average_precision: 0.115926
classification.brier: 0.0525668
classification.log_loss: 0.2051

model_name: multimodal
split: validation
split_role: tuning
n: 4624
regression.mae_days: 270.836
regression.rmse_days: 354.14
regression.median_absolute_error_days: 212.433
regression.r2: -0.178467
regression.mean_error_days: 157.712
classification.accuracy: 0.928849
classification.balanced_accuracy: 0.613034
classification.macro_precision: 0.579597
classification.macro_recall: 0.613034
classification.macro_f1: 0.592471
classification.precision: 0.188
classification.recall: 0.271676
classification.f1: 0.222222
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.802266
classification.average_precision: 0.15468
classification.brier: 0.0618927
classification.log_loss: 0.227649

model_name: multimodal_late_fusion
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 344.121
regression.rmse_days: 425.312
regression.median_absolute_error_days: 290.318
regression.r2: -6.30082
regression.mean_error_days: 326.616
classification.accuracy: 0.99296

classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.49648
classification.macro_f1: 0.498234
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0313177
classification.log_loss: 0.159265

model_name: multimodal_late_fusion
split: test
split_role: test
n: 8299
regression.mae_days: 288.622
regression.rmse_days: 359.638
regression.median_absolute_error_days: 241.319
regression.r2: -0.301349
regression.mean_error_days: 216.964
classification.accuracy: 0.951078
classification.balanced_accuracy: 0.563984
classification.macro_precision: 0.576922
classification.macro_recall: 0.563984
classification.macro_f1: 0.569705
classification.precision: 0.181818
classification.recall: 0.150376
classification.f1: 0.164609
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.819791
classification.average_precision: 0.117625
classification.brier: 0.0537227
classification.log_loss: 0.218114

model_name: multimodal_late_fusion

split: validation
split_role: tuning
n: 4624
regression.mae_days: 277.988
regression.rmse_days: 348.659
regression.median_absolute_error_days: 231.457
regression.r2: -0.142272
regression.mean_error_days: 193.967
classification.accuracy: 0.936419
classification.balanced_accuracy: 0.597521
classification.macro_precision: 0.584467
classification.macro_recall: 0.597521
classification.macro_f1: 0.590387
classification.precision: 0.199005
classification.recall: 0.231214
classification.f1: 0.213904
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.806167
classification.average_precision: 0.169233
classification.brier: 0.0620106
classification.log_loss: 0.241308

model_name: structured
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 515.197
regression.rmse_days: 577.144
regression.median_absolute_error_days: 486.479
regression.r2: -12.4439
regression.mean_error_days: 506.523
classification.accuracy: 0.971489
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.485744
classification.macro_f1: 0.492769
classification.precision: 0

classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0618482
classification.log_loss: 0.26221

model_name: structured
split: test
split_role: test
n: 8299
regression.mae_days: 421.749
regression.rmse_days: 497.501
regression.median_absolute_error_days: 405.994
regression.r2: -1.49029
regression.mean_error_days: 358.646
classification.accuracy: 0.944933
classification.balanced_accuracy: 0.519008
classification.macro_precision: 0.522358
classification.macro_recall: 0.519008
classification.macro_f1: 0.520437
classification.precision: 0.0755556
classification.recall: 0.0639098
classification.f1: 0.0692464
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.717835
classification.average_precision: 0.0732134
classification.brier: 0.077179
classification.log_loss: 0.304804

model_name: structured
split: validation
split_role: tuning
n: 4624
regression.mae_days: 396.182
regression.rmse_days: 469.346

regression.median_absolute_error_days: 379.68
regression.r2: -1.06992
regression.mean_error_days: 320.225
classification.accuracy: 0.934689
classification.balanced_accuracy: 0.518843
classification.macro_precision: 0.521211
classification.macro_recall: 0.518843
classification.macro_f1: 0.519885
classification.precision: 0.0784314
classification.recall: 0.0693642
classification.f1: 0.0736196
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.67274
classification.average_precision: 0.0698108
classification.brier: 0.0849063
classification.log_loss: 0.320023

model_name: structured_extra_trees
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 511.743
regression.rmse_days: 564.05
regression.median_absolute_error_days: 466.041
regression.r2: -11.8408
regression.mean_error_days: 509.97
classification.accuracy: 0.99296
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.49648
classification.macro_f1: 0.498234
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable

classification.average_precision: unavailable
classification.brier: 0.0583006
classification.log_loss: 0.259447

model_name: structured_extra_trees
split: test
split_role: test
n: 8299
regression.mae_days: 413.024
regression.rmse_days: 474.973
regression.median_absolute_error_days: 407.545
regression.r2: -1.26987
regression.mean_error_days: 361.326
classification.accuracy: 0.96349
classification.balanced_accuracy: 0.510419
classification.macro_precision: 0.552927
classification.macro_recall: 0.510419
classification.macro_f1: 0.512777
classification.precision: 0.137255
classification.recall: 0.0263158
classification.f1: 0.044164
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.72502
classification.average_precision: 0.0736308
classification.brier: 0.075062
classification.log_loss: 0.299303

model_name: structured_extra_trees
split: validation
split_role: tuning
n: 4624
regression.mae_days: 388.838
regression.rmse_days: 456.725
regression.median_absolute_error_days: 379.497
regression.r2: -0.960096
regression.mean_error_days: 323.285
classification.accuracy: 0.955882
classification.balanced_accuracy: 0.507629

classification.macro_precision: 0.532852
classification.macro_recall: 0.507629
classification.macro_f1: 0.50758
classification.precision: 0.102564
classification.recall: 0.0231214
classification.f1: 0.0377358
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.664309
classification.average_precision: 0.0675047
classification.brier: 0.0827916
classification.log_loss: 0.319186

model_name: structured_gradient_boosting
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 465.6
regression.rmse_days: 539.15
regression.median_absolute_error_days: 421.329
regression.r2: -10.7321
regression.mean_error_days: 458.404
classification.accuracy: 0.969377
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.484688
classification.macro_f1: 0.492225
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0541099
classification.log_loss: 0.240631

model_name: structured_gradient_boosting
split: test

split_role: test
n: 8299
regression.mae_days: 387.351
regression.rmse_days: 462.039
regression.median_absolute_error_days: 344.211
regression.r2: -1.14793
regression.mean_error_days: 325.038
classification.accuracy: 0.937221
classification.balanced_accuracy: 0.560461
classification.macro_precision: 0.547877
classification.macro_recall: 0.560461
classification.macro_f1: 0.553133
classification.precision: 0.123894
classification.recall: 0.157895
classification.f1: 0.138843
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.719598
classification.average_precision: 0.0811424
classification.brier: 0.0719473
classification.log_loss: 0.285371

model_name: structured_gradient_boosting
split: validation
split_role: tuning
n: 4624
regression.mae_days: 365.975
regression.rmse_days: 444.784
regression.median_absolute_error_days: 321.326
regression.r2: -0.858944
regression.mean_error_days: 288.87
classification.accuracy: 0.929498
classification.balanced_accuracy: 0.552259
classification.macro_precision: 0.544838
classification.macro_recall: 0.552259
classification.macro_f1: 0.548117
classification.precision: 0.123153
classification.recall: 0.144509

classification.f1: 0.132979
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.691914
classification.average_precision: 0.0832157
classification.brier: 0.0789786
classification.log_loss: 0.304067

model_name: structured_hist_gradient_boosting
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 451.916
regression.rmse_days: 544.591
regression.median_absolute_error_days: 381.398
regression.r2: -10.9701
regression.mean_error_days: 442.698
classification.accuracy: 0.966209
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.483105
classification.macro_f1: 0.491407
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0555449
classification.log_loss: 0.235532

model_name: structured_hist_gradient_boosting
split: test
split_role: test
n: 8299
regression.mae_days: 382.581
regression.rmse_days: 474.661
regression.median_absolute_error_days: 319.302

regression.r2: -1.26689
regression.mean_error_days: 312.063
classification.accuracy: 0.930956
classification.balanced_accuracy: 0.55359
classification.macro_precision: 0.537397
classification.macro_recall: 0.55359
classification.macro_f1: 0.543288
classification.precision: 0.103359
classification.recall: 0.150376
classification.f1: 0.122511
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.716619
classification.average_precision: 0.076043
classification.brier: 0.0750317
classification.log_loss: 0.285999

model_name: structured_hist_gradient_boosting
split: validation
split_role: tuning
n: 4624
regression.mae_days: 364.194
regression.rmse_days: 460.262
regression.median_absolute_error_days: 297.399
regression.r2: -0.990567
regression.mean_error_days: 276.495
classification.accuracy: 0.919983
classification.balanced_accuracy: 0.550094
classification.macro_precision: 0.535409
classification.macro_recall: 0.550094
classification.macro_f1: 0.540651
classification.precision: 0.104418
classification.recall: 0.150289
classification.f1: 0.123223
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.693994
classification.average_precision: 0.0806724

classification.brier: 0.0826949
classification.log_loss: 0.306628

model_name: structured_random_forest
split: calibration
split_role: calibration
n: 2841
regression.mae_days: 449.776
regression.rmse_days: 555.866
regression.median_absolute_error_days: 371.885
regression.r2: -11.4708
regression.mean_error_days: 438.385
classification.accuracy: 0.966209
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.483105
classification.macro_f1: 0.491407
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0572916
classification.log_loss: 0.234521

model_name: structured_random_forest
split: test
split_role: test
n: 8299
regression.mae_days: 383.839
regression.rmse_days: 483.275
regression.median_absolute_error_days: 312.278
regression.r2: -1.34991
regression.mean_error_days: 309.034
classification.accuracy: 0.931076
classification.balanced_accuracy: 0.559104
classification.macro_precision: 0.540746

classification.macro_recall: 0.559104
classification.macro_f1: 0.547407
classification.precision: 0.109694
classification.recall: 0.161654
classification.f1: 0.130699
classification.threshold: 0.5
classification.positive_count: 266
classification.auroc: 0.712609
classification.average_precision: 0.0739289
classification.brier: 0.0770733
classification.log_loss: 0.286295

model_name: structured_random_forest
split: validation
split_role: tuning
n: 4624
regression.mae_days: 370.664
regression.rmse_days: 474.574
regression.median_absolute_error_days: 298.367
regression.r2: -1.11629
regression.mean_error_days: 278.006
classification.accuracy: 0.91782
classification.balanced_accuracy: 0.548971
classification.macro_precision: 0.533355
classification.macro_recall: 0.548971
classification.macro_f1: 0.538633
classification.precision: 0.100386
classification.recall: 0.150289
classification.f1: 0.12037
classification.threshold: 0.5
classification.positive_count: 173
classification.auroc: 0.681417
classification.average_precision: 0.0776404
classification.brier: 0.0860553
classification.log_loss: 0.310951

model_name: text
split: calibration
split_role: calibration

n: 2841
regression.mae_days: 324.627
regression.rmse_days: 415.319
regression.median_absolute_error_days: 266.259
regression.r2: -5.96177
regression.mean_error_days: 287.921
classification.accuracy: 0.988384
classification.balanced_accuracy: unavailable
classification.macro_precision: 0.5
classification.macro_recall: 0.494192
classification.macro_f1: 0.497079
classification.precision: 0
classification.recall: unavailable
classification.f1: 0
classification.threshold: 0.5
classification.positive_count: 0
classification.auroc: unavailable
classification.average_precision: unavailable
classification.brier: 0.0288723
classification.log_loss: 0.139608

model_name: text
split: test
split_role: test
n: 8299
regression.mae_days: 284.06
regression.rmse_days: 362.945
regression.median_absolute_error_days: 232.032
regression.r2: -0.325389
regression.mean_error_days: 185.265
classification.accuracy: 0.942282
classification.balanced_accuracy: 0.579432
classification.macro_precision: 0.567488
classification.macro_recall: 0.579432
classification.macro_f1: 0.572827
classification.precision: 0.161905
classification.recall: 0.191729
classification.f1: 0.175559

classification.threshold: 0.5

classification.positive_count: 266

classification.auroc: 0.806332

classification.average_precision: 0.110739

classification.brier: 0.0536982

classification.log_loss: 0.205782

model_name: text

split: validation

split_role: tuning

n: 4624

regression.mae_days: 275.832

regression.rmse_days: 355.961

regression.median_absolute_error_days: 226.229

regression.r2: -0.190616

regression.mean_error_days: 166.457

classification.accuracy: 0.926038

classification.balanced_accuracy: 0.619907

classification.macro_precision: 0.578815

classification.macro_recall: 0.619907

classification.macro_f1: 0.593704

classification.precision: 0.185874

classification.recall: 0.289017

classification.f1: 0.226244

classification.threshold: 0.5

classification.positive_count: 173

classification.auroc: 0.800419

classification.average_precision: 0.161367

classification.brier: 0.0628555

classification.log_loss: 0.231361

Retrospective descriptive research decision support; no clinical validation or guaranteed future accuracy.

Tuning data support model selection; calibration data fit transforms/intervals; only test data provide the untouched held-out estimate. Test inspection can influence subsequent experiments.

Metrics are recomputed from every stored prediction. Saved aggregate metrics are parsed and hashed but are not treated as evidence of accuracy.

Current registry snapshots may contain post-start updates and leakage. Completed-only and convenience samples can create selection bias; chronological splits do not establish prospective validity.

No regression accuracy percentage is defined. Duration errors use days. Binary decisions use probability >= 0.5; confusion matrices have actual rows and predicted columns.

Undefined values are null, displayed as unavailable. One-class splits do not identify class discrimination. Small held-out samples and imbalance can make estimates unstable; no uncertainty interval is invented.

Reliability uses ten equal-width bins. Raw and post-hoc calibrated probabilities are analyzed separately; calibration-cohort diagnostics can be optimistic.

Duration-by-phase subgroup metrics cannot be recomputed because saved duration prediction rows do not contain verified phase labels; stored subgroup aggregates are not substituted.

Charts show all stored points (overlapping points can obscure sample counts); metrics and bin counts include every row.

Tuning validation metrics are optimistic after hyperparameter and raw model selection; test metrics are held out. Calibration diagnostics are in-sample for the calibration transform.

Bootstrap intervals are conditional on the fitted models and assume independent trials; they exclude model-fitting uncertainty and sponsor/site clustering.

Primary probabilities and classification metrics remain raw. Separate sigmoid-calibrated probabilities/metrics exist only with at least 10 calibration rows and two labels of each class; calibration never selects models.

Split-conformal nominal marginal coverage assumes exchangeability, which chronological shift may violate. Finite-sample ranks beyond calibration size produce unbounded intervals; descriptive test coverage is not a future guarantee.

Training range/category indicators describe extrapolation and cannot estimate an OOD probability or guarantee validity.

Nonlinear structured branches and linear modalities share retained cohorts, but have different capacities. Additional complexity does not guarantee improved performance.

Tree impurity importances summarize global behavior, may favor high-cardinality features, and are neither trial-specific SHAP values nor causal evidence.

Text uses TF-IDF; BioBERT, survival, and SHAP are optional separate capabilities and are not claimed as executed.

Completed-only analysis cannot estimate censored/ongoing trial outcomes without a separate survival cohort.

Per-phase test metrics are exploratory; small or one-class subgroups may be unstable or undefined. Scatter plots retain at most 200 chronological rows; metrics and reliability bins use every retained row.

Research decision support only. Completed-only cohort creates follow-up/selection bias; minimum sample guard is not evidence of scientific validity. RETROSPECTIVE SNAPSHOT: baseline feature availability is unverified; post-start registry updates and narrative text may leak outcome information. Do not claim prospective validation.